

Lab1 : back-propagation

110550095 王皓平

1 Introduction

In this assignment, we implemented a simple neural network to classify a set of 2D points based on their position relative to a decision boundary only using Numpy. The model was trained using gradient descent and different activation functions. This report will show the detail of the implementation and discuss about the performance of the model with different training parameters.

2 Implementation Details

2.1 Sigmoid Function

We use Sigmoid as the activate function of the hidden layers.

$$\sigma(x) = \frac{1}{1 + e^{-x}}, \quad 0 < \sigma(x) < 1, \quad \forall x \in \mathbb{R} \quad (1)$$

The sigmoid function is a widely used activation function in neural networks , particularly in the output layer for binary classification tasks. The function maps any value to a range between 0 and 1.

2.2 Neural Network Architecture

The neural network consists of:

- **Input layer:** 2 neurons (x_1, x_2)
- **Hidden layers:** 2 layers, 4 neurons each
- **Output layer:** 1 neuron

2.3 Backpropagation

Using the chain rule, we compute the gradients for each layer:

$$dA_3 = A_3 - y \quad (2)$$

$$dW_3 = \frac{1}{m} A_2^T (dA_3 \cdot \sigma'(A_3)) \quad (3)$$

$$dW_2 = \frac{1}{m} A_1^T (dA_2 \cdot \sigma'(A_2)) \quad (4)$$

$$dW_1 = \frac{1}{m} X^T (dA_1 \cdot \sigma'(A_1)) \quad (5)$$

where $\sigma'(x) = \sigma(x)(1 - \sigma(x))$.

2.4 Weight Update (Gradient Descent)

The weight matrices are updated using gradient descent:

$$W_i = W_i - \alpha dW_i \quad (6)$$

where: α is the learning rate, dW_i is the gradient of the weight matrix.

3 Experimental Results

3.1 Screenshot and comparison figure

Epoch	0	Loss: 0.249764	Epoch	0	Loss: 0.310752
Epoch	5000	Loss: 0.006943	Epoch	5000	Loss: 0.246487
Epoch	10000	Loss: 0.004522	Epoch	10000	Loss: 0.208476
Epoch	15000	Loss: 0.003242	Epoch	15000	Loss: 0.057058
Epoch	20000	Loss: 0.002395	Epoch	20000	Loss: 0.012353
Epoch	25000	Loss: 0.001813	Epoch	25000	Loss: 0.004461
Epoch	30000	Loss: 0.001405	Epoch	30000	Loss: 0.002455
Epoch	35000	Loss: 0.001113	Epoch	35000	Loss: 0.001638
Epoch	40000	Loss: 0.000900	Epoch	40000	Loss: 0.001210
Epoch	45000	Loss: 0.000741	Epoch	45000	Loss: 0.000952
Epoch	50000	Loss: 0.000621	Epoch	50000	Loss: 0.000780
Epoch	55000	Loss: 0.000527	Epoch	55000	Loss: 0.000659
Epoch	60000	Loss: 0.000454	Epoch	60000	Loss: 0.000569
Epoch	65000	Loss: 0.000395	Epoch	65000	Loss: 0.000500
Epoch	70000	Loss: 0.000347	Epoch	70000	Loss: 0.000445
Epoch	75000	Loss: 0.000308	Epoch	75000	Loss: 0.000401
Epoch	80000	Loss: 0.000275	Epoch	80000	Loss: 0.000364
Epoch	85000	Loss: 0.000248	Epoch	85000	Loss: 0.000334
Epoch	90000	Loss: 0.000225	Epoch	90000	Loss: 0.000308
Epoch	95000	Loss: 0.000205	Epoch	95000	Loss: 0.000285

(a) Linear Data

(b) XOR Data

Figure 1: Comparison of Training: Linear vs. XOR

3.2 Accuracy of Predictions

Iter 80	Ground Truth: 0	Prediction: 0.000000
Iter 81	Ground Truth: 0	Prediction: 0.000000
Iter 82	Ground Truth: 0	Prediction: 0.000000
Iter 83	Ground Truth: 0	Prediction: 0.000005
Iter 84	Ground Truth: 0	Prediction: 0.000000
Iter 85	Ground Truth: 1	Prediction: 0.999998
Iter 86	Ground Truth: 0	Prediction: 0.000000
Iter 87	Ground Truth: 0	Prediction: 0.000000
Iter 88	Ground Truth: 0	Prediction: 0.000000
Iter 89	Ground Truth: 0	Prediction: 0.000000
Iter 90	Ground Truth: 1	Prediction: 1.000000
Iter 91	Ground Truth: 1	Prediction: 1.000000
Iter 92	Ground Truth: 0	Prediction: 0.000000
Iter 93	Ground Truth: 1	Prediction: 0.999998
Iter 94	Ground Truth: 1	Prediction: 1.000000
Iter 95	Ground Truth: 0	Prediction: 0.000000
Iter 96	Ground Truth: 1	Prediction: 1.000000
Iter 97	Ground Truth: 1	Prediction: 1.000000
Iter 98	Ground Truth: 0	Prediction: 0.000000
Iter 99	Ground Truth: 1	Prediction: 1.000000
Test Accuracy: 99.00%		


(a) Linear Data

Iter 0	Ground Truth: 0	Prediction: 0.015459
Iter 1	Ground Truth: 1	Prediction: 0.996463
Iter 2	Ground Truth: 0	Prediction: 0.015357
Iter 3	Ground Truth: 1	Prediction: 0.996380
Iter 4	Ground Truth: 0	Prediction: 0.015284
Iter 5	Ground Truth: 1	Prediction: 0.996073
Iter 6	Ground Truth: 0	Prediction: 0.015247
Iter 7	Ground Truth: 1	Prediction: 0.994657
Iter 8	Ground Truth: 0	Prediction: 0.015251
Iter 9	Ground Truth: 1	Prediction: 0.960788
Iter 10	Ground Truth: 0	Prediction: 0.015299
Iter 11	Ground Truth: 0	Prediction: 0.015395
Iter 12	Ground Truth: 1	Prediction: 0.965762
Iter 13	Ground Truth: 0	Prediction: 0.015543
Iter 14	Ground Truth: 1	Prediction: 0.992516
Iter 15	Ground Truth: 0	Prediction: 0.015744
Iter 16	Ground Truth: 1	Prediction: 0.994235
Iter 17	Ground Truth: 0	Prediction: 0.016002
Iter 18	Ground Truth: 1	Prediction: 0.994612
Iter 19	Ground Truth: 0	Prediction: 0.016320
Iter 20	Ground Truth: 1	Prediction: 0.994688
Test Accuracy: 100.00%		


(b) XOR Data

Figure 2: Comparison of Accuracy: Linear vs. XOR

3.3 Learning Curve




(a) Linear Data




(b) XOR Data

Figure 3: Comparison of Learning Curve: Linear vs. XOR

3.4 Visualization



(a) Linear Data




(b) XOR Data


Figure 4: Visualization Plot: Linear vs. XOR

4 Discussion

4.1 Different Learning Rates




(a) Linear Data




(b) XOR Data

4.2 Different Numbers of Hidden Units




(a) Linear Data




(b) XOR Data

4.3 With and Without Activation Functions



(a) Linear Data



(b) XOR Data

5 Questions

5.1 What is the purpose of activation functions?

Activation functions introduce **non-linearity** into the neural network, allowing it to learn complex patterns beyond simple linear relationships. Without an activation function, a neural network would behave like a linear regression model, limiting its ability to solve complex tasks.

5.2 What might happen if the learning rate is too large or too small?

- **if the learning rate is too large:**
 - The model may fail to converge or oscillate around the minimum.
 - It may skip over the optimal solution, preventing proper learning.
- **If the learning rate is too small:**
 - Training will be very slow, requiring many iterations to reach a good solution.
 - The model may get stuck in a local minimum, failing to generalize well.

5.3 What is the purpose of weights and biases in a neural network?

Weights and biases are the key parameters that allow a neural network to learn from data.

- **Weights (W):**
 - Represent the strength of connections between neurons.
 - Higher weights indicate stronger influence of one neuron on another.
 - Are updated during training using gradient descent.
- **Biases (b):**
 - Allow the activation function to shift, helping the model fit better.
 - Ensure neurons activate even when all inputs are zero.

The neural network learns **optimal weights and biases** during training to minimize the loss function and improve predictions.

6 Extra

6.1 Implement Different Optimizers

To improve training performance, we implemented different optimization algorithms:

- **Stochastic Gradient Descent (SGD):**

$$W = W - \alpha \frac{\partial L}{\partial W} \quad (7)$$

- **Adam (Adaptive Moment Estimation):**


$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t \quad (8)$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 \quad (9)$$


$$W = W - \alpha \frac{m_t}{\sqrt{v_t} + \epsilon} \quad (10)$$

- **RMSprop (Root Mean Square Propagation):**

$$W = W - \frac{\alpha}{\sqrt{v_t} + \epsilon} g_t \quad (11)$$



(a) Linear Data



(b) XOR Data

6.2 Implement Different Activation Functions

To compare the performance of different activation functions, we implemented:

- **Sigmoid:**


$$\sigma(x) = \frac{1}{1 + e^{-x}} \quad (12)$$

- **ReLU (Rectified Linear Unit):**


$$f(x) = \max(0, x) \quad (13)$$

- **Tanh (Hyperbolic Tangent):**

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} \quad (14)$$



(a) Linear Data



(b) XOR Data

6.3 Implement Convolutional Layers